

**HW3: Parallelizing Strassen's Matrix-Multiplication Algorithm via OpenMP**

The product  $C$  of two matrices  $A$  and  $B$  of size  $n \times n$  can be computed in  $O(n^3)$  operations using the *standard algorithm* in which each element of the product is obtained by computing the inner product of a row vector of  $A$  and a column vector of  $B$  by employing  $2n$  operations (additions and multiplications between scalars).

```

for i = 1, ... , n
    for j = 1, ... , n
        for k = 1, ... , n
             $C_{ij} = C_{ij} + A_{ik}B_{kj},$ 

```

Volker Strassen proposed a recursive algorithm to compute the product in  $O(n^{2.8})$  operations, and opened the door to the creation of several algorithms that reduced the complexity further. An interesting implication of this work is that the solution of a linear system of order  $n$  can be obtained in  $O(n^{2.8})$ .

Consider the following partitioning of the matrices into *equal sized* blocks:

$$A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}, B = \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}, C = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}, \quad (1)$$

which leads to the straightforward approach to compute  $C$ :

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{bmatrix}, \quad (2)$$

that requires 8 multiplications and 4 additions among matrices of size  $n/2 \times n/2$ .

Strassen's algorithm computes the product as:

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} M_1 + M_4 - M_5 + M_7 & M_3 + M_5 \\ M_2 + M_4 & M_1 - M_2 + M_3 + M_6 \end{bmatrix}, \quad (3)$$

where

$$\begin{aligned} M_1 &= (A_{11} + A_{22})(B_{11} + B_{22}), & M_2 &= (A_{21} + A_{22})B_{11}, & M_3 &= A_{11}(B_{12} - B_{22}), \\ M_4 &= A_{22}(B_{21} - B_{11}), & M_5 &= (A_{11} + A_{12})B_{22}, \\ M_6 &= (A_{21} - A_{11})(B_{11} + B_{12}), & M_7 &= (A_{12} - A_{22})(B_{21} + B_{22}). \end{aligned} \quad (4)$$

Compared to the algorithm in (2), Strassen's approach requires computing only 7 matrix-multiplications, although the number of matrix-additions increases to 18. Since matrix-additions are  $O(n^2)$  whereas multiplications are  $O(n^3)$ , a recursive algorithm using this strategy results in an asymptotic complexity of  $O(n^{\log_2 7})$  which can be approximated by  $O(n^{2.8})$ .

Strassen's *recursive* algorithm applies the idea outlined in (3) and (4) to compute the multiplications for each  $M_i, i = 1, \dots, 7$  recursively. Furthermore, instead of going all the way down to  $1 \times 1$  matrices, you can terminate the recursion after  $k'$  levels ( $k > k' > 1$ ) when you reach matrices of size  $s \times s$  ( $s = n/2^{k'}$ ), and use the standard algorithm to compute matrix-multiplication among these terminal *leaf* matrices.

You are provided with a C++ program `strassen_omp.cpp` that includes code to compute the product of two matrices using the Strassen's algorithm as well the standard method.

To compile and execute the code, use the commands:

```
module load intel
icpx -o strassen_omp.exe -qopenmp strassen_omp.cpp
./strassen_omp.exe <k> <q>
```

where `<k>` and `<q>` are integer arguments that specify the size of the matrix ( $2^k$ ) and the size of the leaf matrices ( $2^q$ ). The output of a sample run is shown below.

```
./strassen_omp.exe 10 4
Matrix size = 1024, Leaf matrix size = 16, Strassen's (s) = 2.3122 s,
Standard = 3.3629 s, Error = 0
```

Note that in its present form, the code is not parallelized.

1. (75 points) Parallelize the code by inserting OpenMP directives to obtain a parallel implementation of Strassen's recursive algorithm.
2. (25 points) Determine the speedup obtained by your code on a single node of Grace using all available cores for matrixes of size  $2^k$  for  $k = 10, \dots, 14$ . Speedup should be computed as the speed improvement over Strassen's algorithm using only a single thread. Experiment with the size of the leaf matrix  $2^q$  to determine which size(s) give you the maximum speedup. Summarize your findings in a document that includes speedup and efficiency graphs as well as your insights into the results you have obtained. Lastly, include a brief description of how to compile and execute the code on Grace.

### Submission:

Upload two files to Canvas:

1. A **single zip file** consisting of the code you developed.
2. Submit a single PDF or MSWord document with your response for Problem.

### Helpful Information:

1. You may use Grace for this assignment.
2. Load the compiler module prior to compiling your program. Use:  
`module load intel`
3. Compile C++ programs with OpenMP pragmas using `icpx` with the switch `-qopenmp`. For example, to compile `code.cpp` to create the executable `code.exe`, use  
`icpx -qopenmp -o code.exe code.cpp`
4. The run time of a code should be measured when it is executed in dedicated mode. Create a batch file and submit your code for execution via the batch system on the system.